

```
In [139]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
sns.set(context="notebook", style = 'darkgrid' , color_codes=True)
from scipy import stats
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
import warnings
warnings.filterwarnings('ignore')
import statsmodels.api as sm
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error, mean_squared_error, mean_absolute_percentage_error
```

In [ ]:

```
In [140]: df=pd.read_csv('https://d2beiqkhq929f0.cloudfront.net/public_assets/assets/000/001/839/original/Jamboree_Admission.csv')
df.head()
```

Out[140]:

	Serial No.	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	Chance of Admit
0	1	337	118	4	4.5	4.5	9.65	1	0.92
1	2	324	107	4	4.0	4.5	8.87	1	0.76
2	3	316	104	3	3.0	3.5	8.00	1	0.72
3	4	322	110	3	3.5	2.5	8.67	1	0.80
4	5	314	103	2	2.0	3.0	8.21	0	0.65


In [141]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Serial No.            500 non-null   int64
1   GRE Score             500 non-null   int64
2   TOEFL Score           500 non-null   int64
3   University Rating     500 non-null   int64
4   SOP                   500 non-null   float64
5   LOR                   500 non-null   float64
6   CGPA                  500 non-null   float64
7   Research              500 non-null   int64
8   Chance of Admit       500 non-null   float64
dtypes: float64(4), int64(5)
memory usage: 35.3 KB
```

## Univariate Analysis


```
In [142]: sns.boxplot(df['GRE Score'])
```

```
Out[142]: <matplotlib.axes._subplots.AxesSubplot at 0x1feb5874940>
```




```
In [143]: sns.boxplot(df['Chance of Admit '])
```

```
Out[143]: <matplotlib.axes._subplots.AxesSubplot at 0x1feb5999ac0>
```




```
In [144]: sns.histplot(df['TOEFL Score'])  
plt.title('Distribution of TOEFL Score')
```

Out[144]: Text(0.5, 1.0, 'Distribution of TOEFL Score')



```
In [145]: sns.countplot(x='University Rating',data=df)  
plt.title('Distribution of University Rating')
```

Out[145]: Text(0.5, 1.0, 'Distribution of University Rating')



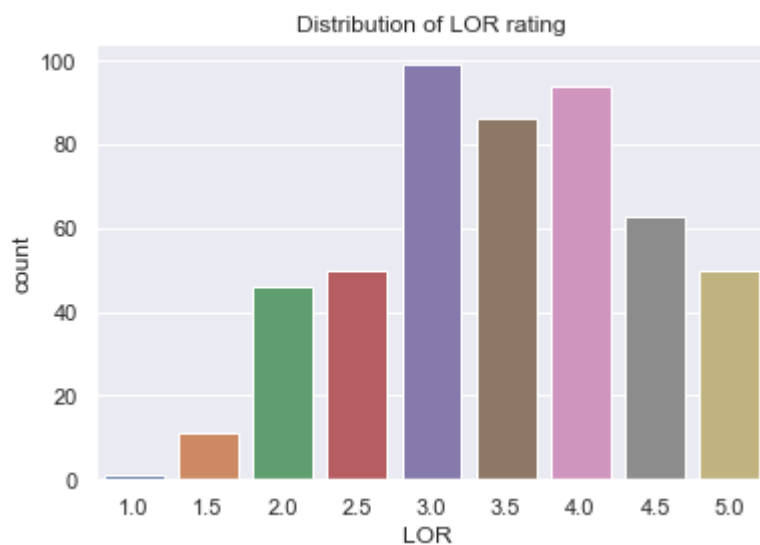
```
In [146]: sns.countplot(x='SOP',data=df)
plt.title('Distribution of SOP Rating')
```

Out[146]: Text(0.5, 1.0, 'Distribution of SOP Rating')




```
In [147]: sns.countplot(x='LOR ',data=df)
plt.title('Distribution of LOR rating')
```

Out[147]: Text(0.5, 1.0, 'Distribution of LOR rating')




```
In [148]: sns.histplot(df['CGPA'])  
plt.title('Distribution of CGPA')
```

Out[148]: Text(0.5, 1.0, 'Distribution of CGPA')



```
In [149]: sns.countplot(x='Research',data=df)  
plt.title('Distribution of Research')
```

Out[149]: Text(0.5, 1.0, 'Distribution of Research')

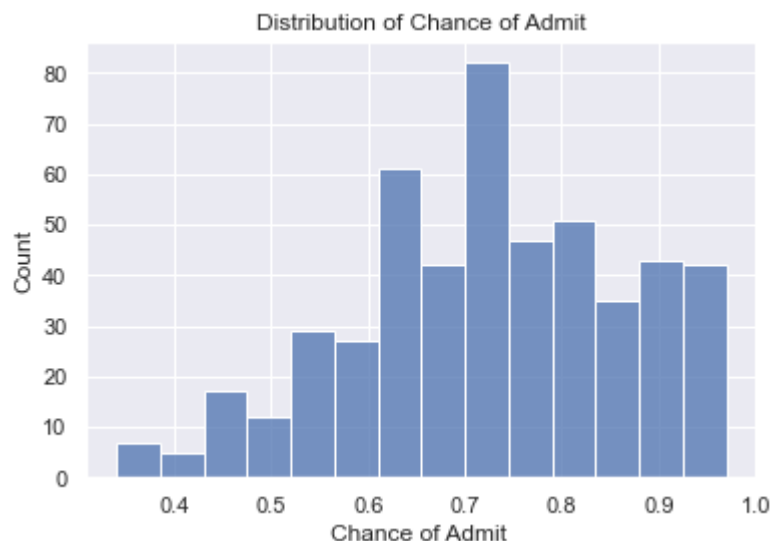


```
In [150]: df['Research'].value_counts()
```

Out[150]: 1 280  
0 220  
Name: Research, dtype: int64

```
In [151]: sns.histplot(df['Chance of Admit '])  
plt.title('Distribution of Chance of Admit')
```


```
Out[151]: Text(0.5, 1.0, 'Distribution of Chance of Admit')
```



## Bivariate Analysis


```
In [152]: sns.scatterplot(x='GRE Score',y='Chance of Admit ',data=df)  
plt.title('Relation between GRE Score and chance of Admit')
```

```
Out[152]: Text(0.5, 1.0, 'Relation between GRE Score and chance of Admit')
```



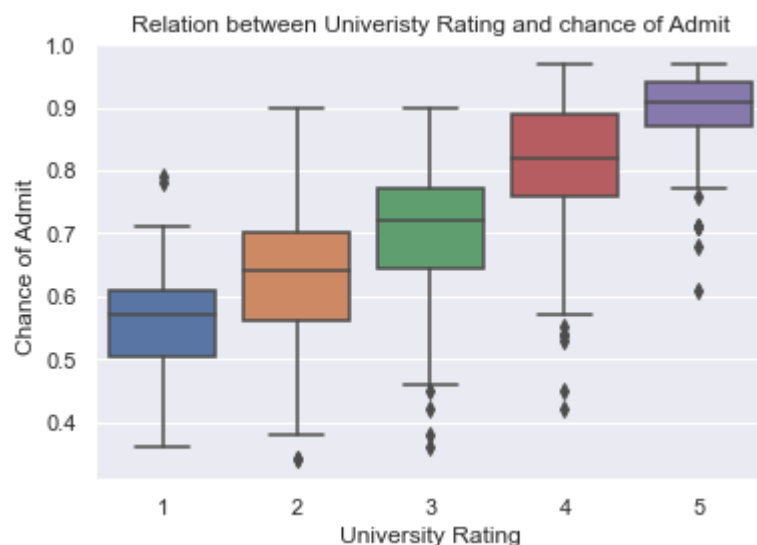
```
In [153]: sns.scatterplot(x='TOEFL Score',y='Chance of Admit ',data=df)
plt.title('Relation between TOEFL score and chance of Admit')
```

Out[153]: Text(0.5, 1.0, 'Relation between TOEFL score and chance of Admit')



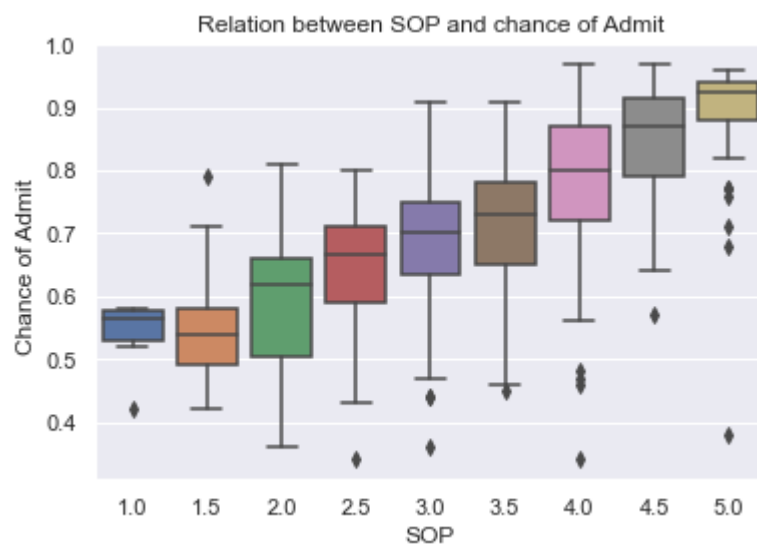
```
In [154]: sns.boxplot(x='University Rating',y='Chance of Admit ',data=df)
plt.title('Relation between Univeristy Rating and chance of Admit')
```

Out[154]: Text(0.5, 1.0, 'Relation between Univeristy Rating and chance of Admit')



```
In [156]: sns.boxplot(x='SOP',y='Chance of Admit ',data=df)
plt.title('Relation between SOP and chance of Admit')
```

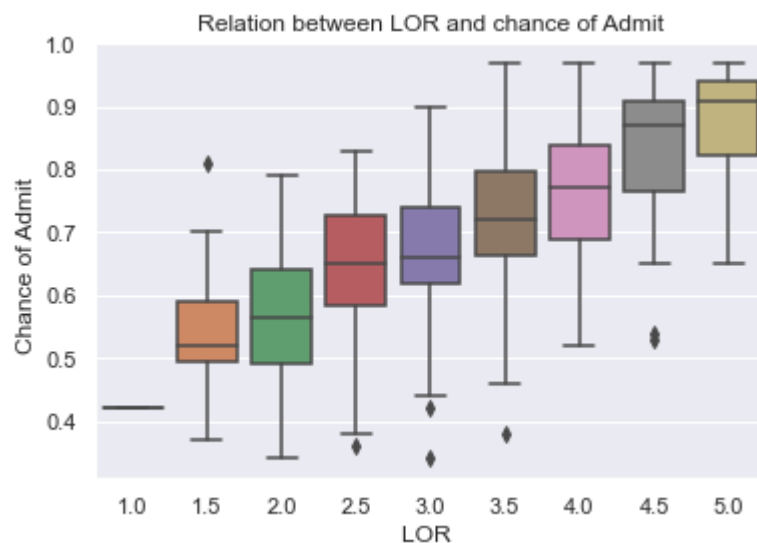
Out[156]: Text(0.5, 1.0, 'Relation between SOP and chance of Admit')



In [ ]:

```
In [157]: sns.boxplot(x='LOR ',y='Chance of Admit ',data=df)
plt.title('Relation between LOR and chance of Admit')
```


Out[157]: Text(0.5, 1.0, 'Relation between LOR and chance of Admit')






```
In [158]: sns.scatterplot(x='CGPA',y='Chance of Admit ',data=df)
plt.title('Relation between CGPA and chance of Admit')
```

Out[158]: Text(0.5, 1.0, 'Relation between CGPA and chance of Admit')



```
In [159]: sns.boxplot(x='Research',y='Chance of Admit ',data=df)
plt.title('Relation between Research and chance of Admit')
```

Out[159]: Text(0.5, 1.0, 'Relation between Research and chance of Admit')



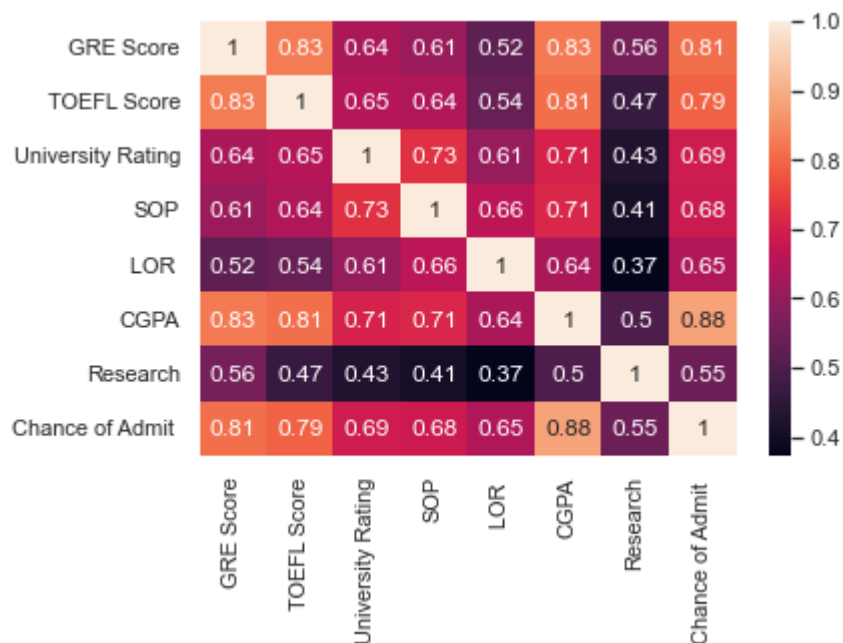
```
In [160]: df.drop(columns=['Serial No.'],axis=1,inplace=True)
df.head()
```

Out[160]:

	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	Chance of Admit
0	337	118	4	4.5	4.5	9.65	1	0.92
1	324	107	4	4.0	4.5	8.87	1	0.76
2	316	104	3	3.0	3.5	8.00	1	0.72
3	322	110	3	3.5	2.5	8.67	1	0.80
4	314	103	2	2.0	3.0	8.21	0	0.65

```
In [161]: sns.heatmap(df.corr(),annot=True)
```

Out[161]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1feb6362640>



```
In [162]: df.shape
```

Out[162]: (500, 8)

```
In [163]: df.columns
```

Out[163]: Index(['GRE Score', 'TOEFL Score', 'University Rating', 'SOP', 'LOR ', 'CGP A',  
'Research', 'Chance of Admit '],  
dtype='object')

Through Trail and error, decided these are the columns that help me in uilding a better model, as there is no much difference in the model with all variables and model with the below mentioned variables

```
In [164]: col=['GRE Score',
'TOEFL Score',
'LOR ',
'CGPA',
'SOP',
'Research'
]
```

## Scaling the data

```
In [165]: from sklearn.preprocessing import StandardScaler
df1=df
for i in col:
    df1[i] = StandardScaler().fit_transform(df1[[i]])
```

```
In [166]: df1.head()
```

Out[166]:

	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	Chance of Admit
0	1.819238	1.778865	4	1.137360	1.098944	1.776806	0.886405	0.92
1	0.667148	-0.031601	4	0.632315	1.098944	0.485859	0.886405	0.76
2	-0.041830	-0.525364	3	-0.377773	0.017306	-0.954043	0.886405	0.72
3	0.489904	0.462163	3	0.127271	-1.064332	0.154847	0.886405	0.80
4	-0.219074	-0.689952	2	-1.387862	-0.523513	-0.606480	-1.128152	0.65

```
In [167]: X=df1[col]
Y=df1['Chance of Admit ']
```

```
In [ ]:
```

```
In [168]: print(X.shape, Y.shape)

(500, 6) (500,)
```

## Performing Linear Regression with Statsmodel

```
In [169]: import statsmodels.api as sm
```

```
In [170]: X_sm = sm.add_constant(X)

sm_model = sm.OLS(Y, X_sm).fit()
```

## Summary of Stats Model

```
In [171]: print(sm_model.summary())
```

# OLS Regression Results

```

=====
Dep. Variable:      Chance of Admit      R-squared:          0.82
1
Model:              OLS      Adj. R-squared:        0.81
9
Method:             Least Squares      F-statistic:        376.
9
Date:               Tue, 15 Mar 2022      Prob (F-statistic):  1.37e-18
0
Time:               23:29:46      Log-Likelihood:     700.1
5
No. Observations:   500      AIC:               -138
6.
Df Residuals:       493      BIC:               -135
7.
Df Model:           6
Covariance Type:    nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.97
5]						
---						
const	0.7217	0.003	268.647	0.000	0.716	0.7
27						
GRE Score	0.0214	0.006	3.775	0.000	0.010	0.0
33						
TOEFL Score	0.0176	0.005	3.329	0.001	0.007	0.0
28						
LOR	0.0164	0.004	4.333	0.000	0.009	0.0
24						
CGPA	0.0728	0.006	12.531	0.000	0.061	0.0
84						
SOP	0.0042	0.004	0.991	0.322	-0.004	0.0
12						
Research	0.0123	0.003	3.767	0.000	0.006	0.0
19						

```

=====
Omnibus:           112.527      Durbin-Watson:      0.78
9
Prob(Omnibus):     0.000      Jarque-Bera (JB):   260.84
9
Skew:              -1.158      Prob(JB):           2.28e-5
7
Kurtosis:          5.675      Cond. No.            5.2
5
=====

```

## Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

## Checking VIF Score

```
In [172]: from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```
In [173]: vif = pd.DataFrame()
X_t = X
vif['Features'] = X_t.columns
vif['VIF'] = [variance_inflation_factor(X_t.values, i) for i in range(X_t.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

Out[173]:

	Features	VIF
3	CGPA	4.68
0	GRE Score	4.45
1	TOEFL Score	3.87
4	SOP	2.45
2	LOR	1.99
5	Research	1.49

## Linear Regression Using Scikit Learn

```
In [174]: from sklearn.linear_model import LinearRegression
rm=LinearRegression()
rm.fit(X,Y)
```

Out[174]: LinearRegression()

```
In [175]: print(rm.intercept_)
```

0.7217399999999997

```
In [189]: [print(f"The coefficient of {col[i]} is {rm.coef_[i]}") for i in range(len(col))]
```

The coefficient of GRE Score is 0.021404695703631405  
The coefficient of TOEFL Score is 0.01760495005366209  
The coefficient of LOR is 0.016432612597028093  
The coefficient of CGPA is 0.07283664907806046  
The coefficient of SOP is 0.004168721982917806  
The coefficient of Research is 0.012349581379967556

Out[189]: [None, None, None, None, None, None]

```
In [177]: rm.score(X,Y)
```

```
Out[177]: 0.8210166861582956
```


## Mean of Residuals

```
In [178]: Y_pred = rm.predict(X)
residuals = Y.values-Y_pred
mean_residuals = np.mean(residuals)
print("Mean of Residuals {}".format(round(mean_residuals,5)))
```

```
Mean of Residuals 0.0
```

## Test of Homoscedasticity

```
In [179]: sns.scatterplot(Y_pred,residuals)
plt.xlabel('y_pred/predicted values')
plt.ylabel('Residuals')
plt.xlabel('y_pred/predicted values')
plt.ylabel('Residuals')
plt.ylim(-1,1)
plt.xlim(0.3,1.10)
p = sns.lineplot([0.37,1.05],[0,0],color='green')
```



## Goldfeld Quandt Test to check Homoscedasticity


```
In [180]: import statsmodels.stats.api as sms
from statsmodels.compat import lzip
name=['F_stat','P_value']
GQT = sms.het_goldfeldquandt(residuals, X)
lzip(name,GQT)
```

```
Out[180]: [('F_stat', 0.45996749764483563), ('P_value', 0.9999999989223426)]
```

## Normality of Residuals

```
In [181]: fig, axes = plt.subplots(2)
sns.distplot(residuals, ax=axes[0])
sm.qqplot(residuals, ax=axes[1])
p = stats.normaltest(residuals).pvalue
if p > 0.05:
    print('Residuals are Normally Distributed as p value is ' + str(p) + ' which is greater than 0.05')
else:
    print('Residuals are not Normally Distributed as p value is ' + str(p) + ' which is less than 0.05')
```

Residuals are not Normally Distributed as p value is 3.6739042749043505e-25 which is less than 0.05



## Performance checks on Model

```
In [182]: print("mean absolute error : ", mean_absolute_error(Y_pred, Y))
print("Mean Squared error : ", mean_squared_error(Y_pred, Y))
print("Root Mean squared error : ", np.sqrt(mean_squared_error(Y_pred, Y)))
print("Mean absolute Percentage error : ", mean_absolute_percentage_error(Y_pred, Y))
print("R-squared : ", rm.score(X, Y))
print("Adjusted R-squared : ", 1 - (1 - rm.score(X, Y)) * (len(Y) - 1) / (len(Y) - X.shape[1] - 1))
```

mean absolute error : 0.04276860707461887  
Mean Squared error : 0.003558326525884695  
Root Mean squared error : 0.059651710167309496  
Mean absolute Percentage error : 0.0641908511172667  
R-squared : 0.8210166861582956  
Adjusted R-squared : 0.8188383902494716



## Linear Regression with Train and Test data

```
In [183]: x_train,x_test,y_train,y_test = train_test_split(X,Y,test_size = 0.2,random_state=1 )
```

```
In [184]: fm=LinearRegression()  
fm.fit(x_train,y_train)
```

```
Out[184]: LinearRegression()
```

```
In [185]: print(fm.intercept_,fm.coef_)  
  
0.7226452460475254 [0.02142235 0.01981272 0.01430457 0.07165618 0.00570992 0.01020389]
```

```
In [186]: y_pred = fm.predict(x_test)
```

```
In [187]: print("mean absolute error : ",mean_absolute_error(y_pred,y_test))  
print("Mean Squared error : ", mean_squared_error(y_pred,y_test))  
print("Root Mean squared error : ", np.sqrt(mean_squared_error(y_pred,y_test)))  
print("Mean absolute Percentage error : ",mean_absolute_percentage_error(y_pred,y_test))  
print("R-squared : ",rm.score(x_train,y_train))  
print("Adjusted R-squared : ", 1 - (1-rm.score(x_train, y_train))*(len(y_train)-1)/(len(y_train)-x_train.shape[1]-1))  
  
mean absolute error : 0.04006870393111951  
Mean Squared error : 0.0034716563533543753  
Root Mean squared error : 0.058920763346670714  
Mean absolute Percentage error : 0.05874548625702469  
R-squared : 0.8202210391178049  
Adjusted R-squared : 0.8174763221577713
```

```
In [190]: [print(f"The coefficient of {col[i]} is {fm.coef_[i]}") for i in range(len(col))]  
  
The coefficient of GRE Score is 0.021422347634471903  
The coefficient of TOEFL Score is 0.019812718838837576  
The coefficient of LOR is 0.014304568752076963  
The coefficient of CGPA is 0.07165617630233034  
The coefficient of SOP is 0.005709916757106009  
The coefficient of Research is 0.010203891926236885
```

```
Out[190]: [None, None, None, None, None, None]
```

For all the models which are built using different techniques, there is a good value of R-squared and adjusted R-Square value. So we can consider our model as good and the Root Mean Squared Errors for Scikit Learn Models (one is with direct data and other with train, test data) are also pretty low